

## Week 3: Summarizing Data II

Wenhao Jiang

Department of Sociology  
New York University

September 23, 2022

# Data

# Data Structures

- ▶ R also has a number of basic data *structures*.
- ▶ A data structure is either
  - ▶ homogeneous (all elements are of the same data type)
  - ▶ heterogeneous (elements can be of more than one data type).

Dimension	Homogeneous	Heterogeneous
1	Vector	List
2	Matrix	Data Frame
3+	Array	

## Summarize Data (Basics)

## Read Data

- ▶ A **data frame** is the **most** common way that we store and interact with data

```
## set working directory
```

```
setwd("~/Dropbox/Teaching/SOCUA-302/Week 2")
```

```
## read the file
```

```
gss <- read.csv("GSS_SOCUA_W2.csv")
```

## Combine data subsetting and summarizing (base R)

- ▶ Some values in the variable `sibs` are not real observations
- ▶ Negative values are missingness
- ▶ How do we calculate the mean value of `sibs` without negative `sibs` values?

```
mean(gss[gss$sibs>=0,"sibs"])
```

```
## [1] 3.863355
```

## Exercise

- ▶ How do we calculate the variance of sibs in year 2018 only?

## Combine data subsetting and summarizing (dplyr)

- ▶ We can also use dplyr to subset and summarize data
  - ▶ Remember `filter()` filters rows and `select()` select columns
- ▶ How do we calculate the mean value of `sibs` without negative `sibs` values?

```
library(dplyr)
gss %>% filter(sibs>0) %>% summarize(sibs_mean = mean(sibs))
```

```
##    sibs_mean
## 1    4.066925
```

- ▶ `summarize()` follows the format: `summarize(your_summarize_name = function(variable))`



## Combine data subsetting and summarizing (dplyr)

- ▶ How do we calculate the mean value of sibs without negative sibs values?
- ▶ Why bothering using dplyr and summarize?
  - ▶ Because we can summarize data using different functions at the same time

```
gss %>% filter(sibs>0) %>% summarize(sibs_mean = mean(sibs),  
                                     sibs_var = var(sibs))
```

```
##      sibs_mean sibs_var  
## 1    4.066925 9.620029
```

## Export the data

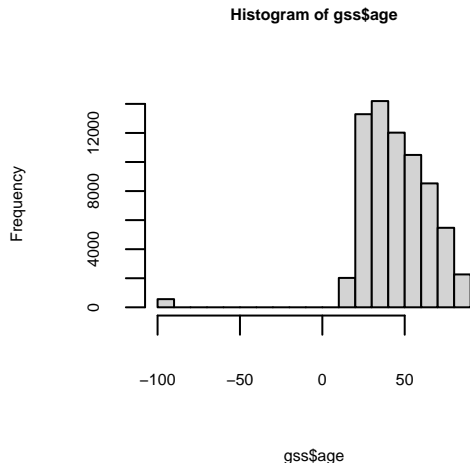
- ▶ `write.csv` save (or export) the dataframe in `.csv` format.

```
write.csv(gss, "cleaned_gss.csv")
```

# Histogram plot

- Histograms describe the distribution of **numeric** data


```
hist(gss$age)
```



- ▶ Again, we want to drop negative values of age
- ▶ We may also want to see the distribution for some years but not others

```
hist(gss[gss$age>0 & gss$year<=2018 & gss$year>=2012,"age"])
```

Histogram of gss[gss\$age > 0 & gss\$year <= 2018 & gss\$year >= 2012, '




```
gss[gss$age > 0 & gss$year <= 2018 & gss$year >= 2012, "age"]
```

- To change plot titles and x-axis label

```
hist(gss[gss$age>0 & gss$year<=2018 & gss$year>=2012,"age"],  
     main = "Distribution of Respondents Ages in 2012-2018, GSS",  
     xlab = "Age")
```


Distribution of Respondents Ages in 2012-2018, GSS



► To control the number of bins

```
hist(gss[gss$age>0 & gss$year<=2018 & gss$year>=2012,"age"],  
     main = "Distribution of Respondents Ages in 2012-2018, GSS",  
     xlab = "Age",  
     breaks = 30)
```

Distribution of Respondents Ages in 2012-2018, GSS



# Barplot


- ▶ Barplots describe the distribution of **categorical** data

```
## create a count summary in each category by function `table`  
counts <- table(gss[gss$age>0 & gss$year<=2018 &  
                gss$year>=2012 & gss$marital>=0, "marital"])  
counts
```

```
##  
##      1      2      3      4      5  
## 4250  821 1622  325 2673
```

```
barplot(counts, main="Distribution of Marital Status in 2012-2018, GSS",  
        xlab="Marital Status",  
        names.arg=c("Married", "Widowed", "Divorced", "Separated",  
                    "Never Married"),  
        cex.lab=0.5, cex.axis=0.5, cex.main=0.5, cex.sub=0.5, cex.names=0.32)
```

Distribution of Marital Status in 2012-2018, GSS



Marital Status




- ▶ We can also change the order of the categories by first factorizing the variable with manual levels

```
gss$marital <- factor(gss$marital, levels = c(1,3,4,2,5))
counts <- table(gss[gss$age>0 & gss$year<=2018 &
  gss$year>=2012, "marital"])
```

```
barplot(counts, main="Distribution of Marital Status in 2012-2018, GSS",  
        xlab="Marital Status",  
        names.arg=c("Married", "Divorced", "Separated", "Widowed", "Never Married"),  
        cex.lab=0.5, cex.axis=0.5, cex.main=0.5, cex.sub=0.5, cex.names=0.35)
```

Distribution of Marital Status in 2012-2018, GSS



## Summarize Data (Advanced)

## Summarize data by group

- ▶ In many cases, we not only want data summaries for the total sample, but summaries for some sub-samples
- ▶ For example, we may want to summarize religious preference by gender (the column `relig`)
- ▶ In the `relig` column, 4 means no religious belief, and other non-zero numbers mean some religious belief
- ▶ It is a non-ordered categorical variable

```
unique(gss$relig)
```

```
## [1] 3 2 1 5 4 -99 -98 13 11 9 6 10 12 7 8 -97
```

## Summarize data by group

- ▶ As there are many religious categories, we may only want two groups categorization for now: religious and non-religious individuals

```
gss[gss$relig!=4 & gss$relig>0,  
     "relig"] <- 1  
gss[gss$relig==4,  
     "relig"] <- 0  
unique(gss$relig)
```

```
## [1] 1 0 -99 -98 -97
```

## Summarize data by group

- ▶ We will use the `group_by()` function in `dplyr` to summarize data by the group we specify
- ▶ It can be nested in a `%>%` pipeline

```
gss %>%  
  filter(relig>=0&sex>0) %>%  
  group_by(sex) %>%  
  summarize(religious = mean(relig))
```

```
## # A tibble: 2 x 2  
##   sex religious  
##   <int>     <dbl>  
## 1     1     0.835  
## 2     2     0.897
```

## Exercise

- ▶ The calculation above includes all individuals surveyed since 1972 until 2021
- ▶ How to get the same statistic, but only in 2021?

## Exercise

- ▶ The variable `relig16` asks which religion was the respondent raised in
- ▶ How to get the proportion of religious people (at the time of being interviewed, the column `relig`) by whether they were raised in a religious family?
- ▶ Do not forget to drop missing values first!



## Exercise

```
gss[gss$relig16!=4&gss$relig16>=0,"relig16"] <- 1
gss[gss$relig16==4,"relig16"] <- 0
gss %>%
  filter(relig>=0&sex>0&relig16>=0) %>%
  group_by(relig16) %>%
  summarize(religious = mean(relig))
```

```
## # A tibble: 2 x 2
##   relig16 religious
##   <dbl>      <dbl>
## 1       0      0.450
## 2       1      0.904
```

## Exercise

- ▶ The variable `relig16` asks which religion was the respondent raised in
- ▶ How to get the proportion of religious people (at the time of being interviewed, the column `relig`) by whether they were raised in a religious family?
- ▶ How to add another column summarizing the proportion of non-religious people by whether they were raised in a religious family?

## Exercise

```
gss[gss$relig16!=4&gss$relig16>=0,"relig16"] <- 1
gss[gss$relig16==4,"relig16"] <- 0
gss %>%
  filter(relig>=0&sex>0&relig16>=0) %>%
  group_by(relig16) %>%
  summarize(religious = mean(relig),
            nonreligious=mean(1-relig))
```

```
## # A tibble: 1 x 3
##   relig16 religious nonreligious
##   <dbl>      <dbl>      <dbl>
## 1         1      0.879      0.121
```

## Visualizing the trend

- ▶ We have already calculated the proportion of religious men and women in all and one specific year(s)
- ▶ Another common exploration of such data is to analyze the temporal trend of the proportion of religious people by gender over the years

```
gss %>%  
  filter(relig>=0&sex>0) %>%  
  group_by(year, sex) %>%  
  summarize(religious = mean(relig)) %>%  
  head()
```

```
## # A tibble: 6 x 3  
## # Groups:   year [3]  
##   year    sex religious  
##   <int> <int>     <dbl>  
## 1  1972     1     0.932  
## 2  1972     2     0.965  
## 3  1973     1     0.911  
## 4  1973     2     0.958  
## 5  1974     1     0.903  
## 6  1974     2     0.957
```

## Visualizing the trend (ggplot2)

- ▶ While base R can handle such visualizations, we will introduce another powerful visualization package, `ggplot2`
- ▶ `ggplot2` can also be nested in the `%>%` pipeline
- ▶ It is highly flexible in changing colors, setting fonts, customizing legends, etc.

## Viasualizing the trend (ggplot2)

- ▶ The basic structure of ggplot2 is


```
## create a ggplot object  
ggplot(data,  
        aes(x=variable1,  
            y=variable2,  
            group=variable3,  
            ...)) +  
  geom_line() +  
  geom_point() +  
  ...
```

- ▶ The basic **parameters** such as x-axis, y-axis, groups and colors, are controlled by `aes()` that stands for aesthetics
- ▶ To create a trend plot, we need a **line** plot controlled by `geom_line()`
- ▶ There are many other options, including `geom_point()` that creates scatter plot, `geom_histogram()` that creates histograms, etc.
- ▶ Check this website for much more details of what `ggplot2` can do!
- ▶ <https://ggplot2.tidyverse.org/reference/>



## Visualizing the trend (ggplot2)


```
library(ggplot2)
gss %>%
  filter(relig>=0&sex>0) %>%
  group_by(year, sex) %>%
  summarize(religious = mean(relig)) %>%
  ggplot(aes(x=year, y=religious, group=sex, color=sex)) +
  geom_line()
```



- ▶ It looks fine, but need some refinements!


- ▶ It looks fine, but need some refinements!
- ▶ R identifies the variable `sex` as a continuous variable, we need a categorical variable
  - ▶ We can change the data type of `sex` into characters
  - ▶ We use function `mutate(new_variable = ...)` in `dplyr` to create new variable(s)
  - ▶ The new variable name can be the same as the original one (i.e., overwriting)

```
gss %>%  
  filter(relig>=0 & sex>=0) %>%  
  mutate(sex = as.character(sex)) %>%  
  group_by(year, sex) %>%  
  summarize(religious = mean(relig)) %>%  
  ggplot(aes(x=year, y=religious, color=sex)) +  
  geom_line()
```



- ▶ Readers do not know what 1 and 2 in sex means! We need male and female labels
- ▶ We may also want to change the line color

```
gss %>%  
  filter(relig>=0 & sex>=0) %>%  
  mutate(sex = as.character(sex)) %>%  
  group_by(year, sex) %>%  
  summarize(religious = mean(relig)) %>%  
  ggplot(aes(x=year, y=religious, color=sex)) +  
  scale_color_manual(labels = c("Male", "Female"),  
                     values = c("darkblue", "darkred"))  
  geom_line()
```




- ▶ The background looks a bit pale...
- ▶ And we need a title

```
gss %>%  
  filter(relig>=0 & sex>=0) %>%  
  mutate(sex = as.character(sex)) %>%  
  group_by(year, sex) %>%  
  summarize(religious = mean(relig)) %>%  
  ggplot(aes(x=year,y=religious,color=sex)) +  
  scale_color_manual(labels = c("Male", "Female"),  
                     values = c("darkblue", "darkred")) +  
  geom_line() +  
  theme_bw() +  
  ggtitle("The Proportion of Religious People \n  
          by Gender, 1972-2021, GSS")
```



## The Proportion of Religious People by Gender, 1972–2021, GSS



With more tweaks, we can make the figure publishable

```
gss %>%
  filter(relig>=0 & sex>=0) %>%
  mutate(sex = as.character(sex)) %>%
  group_by(year, sex) %>%
  summarize(religious = mean(relig)) %>%
  ggplot(aes(x=year,y=religious,color=sex)) +
  geom_line() +
  scale_color_manual(labels = c("Male", "Female"),
                     values = c("darkblue", "darkred")) +
  theme_bw() +
  ggtitle("The Proportion of Religious People by Gender, 1972-2021, GSS") +
  xlab("Year") +
  ylab("Proportion with Religious Preferences") +
  theme(plot.title=element_text(size=10),
        text=element_text(family="Times"),
        axis.title.x=element_text(size=8),
        axis.title.y=element_text(size=8),
        legend.title = element_text(size=8),
        legend.text = element_text(size=6))
```

